

```
In [ ]: #Wania Urooj Suleman CMSID:49178
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
```

```
In [2]: def plot_confusion_matrix(y,y_predict):
        "this function plots the confusion matrix"
        from sklearn.metrics import confusion_matrix
        cm = confusion_matrix(y, y_predict)
        ax= plt.subplot()
        sns.heatmap(cm, annot=True, ax = ax);
        ax.set_xlabel('Predicted labels')
        ax.set_ylabel('True labels')
        ax.set_title('Confusion Matrix');
        ax.xaxis.set_ticklabels(['did not land', 'land']); ax.yaxis.set_ticklabels(['did not land', 'landed'])
```

```
In [5]: data = pd.read_csv("https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-DS0321EN-SkillsNetwork/datasets/dataset_part_2.csv")
data.head()
data.to_csv('dataset_part_2-2.csv')
```

```
In [8]: X = pd.read_csv('https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-DS0701EN-SkillsNetwork/api/dataset_part_3.csv')
X.head(100)
```

```
Out[8]:
```

	FlightNumber	PayloadMass	Flights	Block	ReusedCount	Orbit_ES-L1	Orbit_GEO	Orbit_GTO	Orbit_HEO	Orbit_ISS	...	Serial_B1058	Serial_B1059	Serial_B1060	Serial_B1062	GridFins_False	Gric
0	1.0	6104.959412	1.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	1.0	
1	2.0	525.000000	1.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	1.0	
2	3.0	677.000000	1.0	1.0	0.0	0.0	0.0	0.0	0.0	1.0	...	0.0	0.0	0.0	0.0	1.0	
3	4.0	500.000000	1.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	1.0	
4	5.0	3170.000000	1.0	1.0	0.0	0.0	0.0	1.0	0.0	0.0	...	0.0	0.0	0.0	0.0	1.0	
...
85	86.0	15400.000000	2.0	5.0	2.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	1.0	0.0	0.0	
86	87.0	15400.000000	3.0	5.0	2.0	0.0	0.0	0.0	0.0	0.0	...	1.0	0.0	0.0	0.0	0.0	
87	88.0	15400.000000	6.0	5.0	5.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	
88	89.0	15400.000000	3.0	5.0	2.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	1.0	0.0	0.0	
89	90.0	3681.000000	1.0	5.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	1.0	0.0	

90 rows × 83 columns

```
In [9]: Y = data['Class'].to_numpy()
```

```
In [10]: transform = preprocessing.StandardScaler()  
X = transform.fit_transform(X)
```

```
In [11]: X
```

```
Out[11]: array([[ -1.71291154e+00,  -1.94814463e-16,  -6.53912840e-01, ...,  
                -8.35531692e-01,   1.93309133e+00,  -1.93309133e+00],  
               [ -1.67441914e+00,  -1.19523159e+00,  -6.53912840e-01, ...,  
                -8.35531692e-01,   1.93309133e+00,  -1.93309133e+00],  
               [ -1.63592675e+00,  -1.16267307e+00,  -6.53912840e-01, ...,  
                -8.35531692e-01,   1.93309133e+00,  -1.93309133e+00],  
               ...,  
               [  1.63592675e+00,   1.99100483e+00,   3.49060516e+00, ...,  
                1.19684269e+00,  -5.17306132e-01,   5.17306132e-01],  
               [  1.67441914e+00,   1.99100483e+00,   1.00389436e+00, ...,  
                1.19684269e+00,  -5.17306132e-01,   5.17306132e-01],  
               [  1.71291154e+00,  -5.19213966e-01,  -6.53912840e-01, ...,  
                -8.35531692e-01,  -5.17306132e-01,   5.17306132e-01]])
```

```
In [12]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=2)
```

```
In [13]: Y_test.shape
```

```
Out[13]: (18,)
```

```
In [14]: parameters = {'C':[0.01,0.1,1],  
                       'penalty':['l2'],  
                       'solver':['lbfgs']}
```

```
In [15]: parameters = {"C":[0.01,0.1,1], 'penalty':['l2'], 'solver':['lbfgs']}  
lr=LogisticRegression()  
logreg_cv = GridSearchCV(lr, parameters, cv = 10)  
logreg_cv.fit(X_train, Y_train)
```

```
Out[15]: GridSearchCV(cv=10, estimator=LogisticRegression(),  
                      param_grid={'C': [0.01, 0.1, 1], 'penalty': ['l2'],  
                                   'solver': ['lbfgs']})
```


```
In [16]: print("tuned hyperparameters :(best parameters) ",logreg_cv.best_params_)  
print("accuracy :",logreg_cv.best_score_)  
  
tuned hyperparameters :(best parameters) {'C': 0.01, 'penalty': 'l2', 'solver': 'lbfgs'}  
accuracy : 0.8464285714285713
```

```
In [17]: method = []  
accuracy = []  
method.append('Logistic regression')  
accuracy.append(logreg_cv.score(X_test, Y_test))
```

```
In [18]: method, accuracy
```

```
Out[18]: (['Logistic regression'], [0.8333333333333334])
```

```
In [19]: yhat=logreg_cv.predict(X_test)
plot_confusion_matrix(Y_test,yhat)
```



```
In [20]: parameters = {'kernel':('linear', 'rbf','poly','rbf', 'sigmoid'),
                        'C': np.logspace(-3, 3, 5),
                        'gamma':np.logspace(-3, 3, 5)}
svm = SVC()
```

```
In [21]: svm_cv = GridSearchCV(svm, parameters, cv = 10)
svm_cv.fit(X_train, Y_train)
```

```
Out[21]: GridSearchCV(cv=10, estimator=SVC(),
                      param_grid={'C': array([1.00000000e-03, 3.16227766e-02, 1.00000000e+00, 3.16227766e+01,
1.00000000e+03]),
                      'gamma': array([1.00000000e-03, 3.16227766e-02, 1.00000000e+00, 3.16227766e+01,
1.00000000e+03]),
                      'kernel': ('linear', 'rbf', 'poly', 'rbf', 'sigmoid')})
```


```
In [22]: print("tuned hyperparameters :(best parameters) ",svm_cv.best_params_)
print("accuracy :",svm_cv.best_score_)
```

```
tuned hyperparameters :(best parameters) {'C': 1.0, 'gamma': 0.03162277660168379, 'kernel': 'sigmoid'}
accuracy : 0.8482142857142856
```

```
In [23]: method.append('Support vector machine')
accuracy.append(svm_cv.score(X_test, Y_test))
print("test set accuracy :",svm_cv.score(X_test, Y_test))
```

```
test set accuracy : 0.8333333333333334
```

```
In [24]: yhat=svm_cv.predict(X_test)
plot_confusion_matrix(Y_test,yhat)
```



```
In [25]: parameters = {'criterion': ['gini', 'entropy'],
'splitter': ['best', 'random'],
'max_depth': [2*n for n in range(1,10)],
'max_features': ['auto', 'sqrt'],
'min_samples_leaf': [1, 2, 4],
'min_samples_split': [2, 5, 10]}

tree = DecisionTreeClassifier()
```

```
In [26]: tree_cv = GridSearchCV(tree,parameters,cv=10)
tree_cv.fit(X_train, Y_train)
```

```
Out[26]: GridSearchCV(cv=10, estimator=DecisionTreeClassifier(),
param_grid={'criterion': ['gini', 'entropy'],
'max_depth': [2, 4, 6, 8, 10, 12, 14, 16, 18],
'max_features': ['auto', 'sqrt'],
'min_samples_leaf': [1, 2, 4],
'min_samples_split': [2, 5, 10],
'splitter': ['best', 'random']})
```


```
In [27]: print("tuned hyperparameters :(best parameters) ",tree_cv.best_params_)
print("accuracy :",tree_cv.best_score_)

tuned hyperparameters :(best parameters) {'criterion': 'entropy', 'max_depth': 4, 'max_features': 'sqrt', 'min_samples_leaf': 2, 'min_samples_split': 2, 'split
ter': 'best'}
accuracy : 0.8892857142857145
```

```
In [28]: method.append('Decision tree classifier')
accuracy.append(tree_cv.score(X_test, Y_test))
print("test set accuracy :",tree_cv.score(X_test, Y_test))

test set accuracy : 0.9444444444444444
```

```
In [29]: yhat = svm_cv.predict(X_test)
plot_confusion_matrix(Y_test,yhat)
```



```
In [30]: parameters = {'n_neighbors': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
                        'algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute'],
                        'p': [1,2]}

KNN = KNeighborsClassifier()
```

```
In [31]: knn_cv = GridSearchCV(KNN,parameters,cv=10)
knn_cv.fit(X_train, Y_train)
```

```
Out[31]: GridSearchCV(cv=10, estimator=KNeighborsClassifier(),
                      param_grid={'algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute'],
                                   'n_neighbors': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
                                   'p': [1, 2]})
```


```
In [32]: print("tuned hpyerparameters :(best parameters) ",knn_cv.best_params_)
print("accuracy :",knn_cv.best_score_)

tuned hpyerparameters :(best parameters) {'algorithm': 'auto', 'n_neighbors': 10, 'p': 1}
accuracy : 0.8482142857142858
```

```
In [33]: method.append('K nearest neighbors')
accuracy.append(knn_cv.score(X_test, Y_test))
print("test set accuracy :",knn_cv.score(X_test, Y_test))

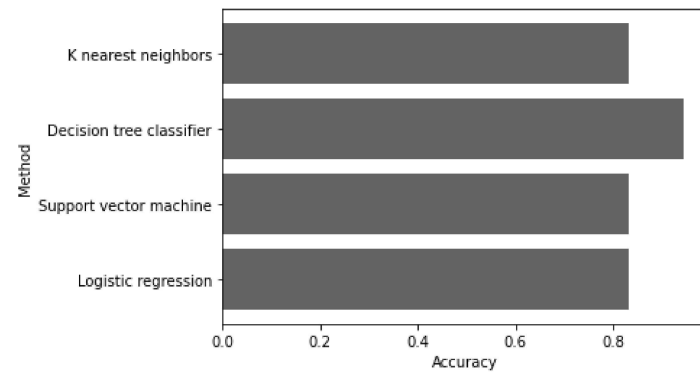
test set accuracy : 0.8333333333333334
```

```
In [34]: yhat = knn_cv.predict(X_test)
plot_confusion_matrix(Y_test,yhat)
```



```
In [35]: import numpy as np
import matplotlib.pyplot as plt

plt.barh(method, accuracy)
plt.xlabel('Accuracy')
plt.ylabel('Method')
plt.show()
```



```
In [36]: results_df = {'method': method,
'accuracy': accuracy}

frame = pd.DataFrame(results_df)
frame
```

```
Out[36]:
```

	method	accuracy
0	Logistic regression	0.833333
1	Support vector machine	0.833333
2	Decision tree classifier	0.944444
3	K nearest neighbors	0.833333

